


## CO513: Advanced Computer Communication Networks - Lab 02

E/19/446: Wijerathna I.M.K.D.I.

09/05/2025

### Network Configuration and Setup Summary



A client machine was added to the setup for a clearer and more realistic simulation. Instead of using iperf3, I used Python HTTP servers and clients due to issues with iperf3's dual-connection design (data and control) and single-client limitation, which made consistent load balancing difficult. Since, this approach consists of a single-connection design it provided better control and reliability during testing.

Docker and Docker Compose were used to simulate these nodes within the same virtual network. Please refer the [Appendix 1](#) and [2](#) for the Docker configurations.

```
[~/C0513/Lab-2]
kalindu ➤ sudo docker ps
[sudo] password for kalindu:
CONTAINER ID   IMAGE          COMMAND                  CREATED        STATUS        PORTS          NAMES
6c3da392631c   lab-2-client   "bash"                  About a minute ago Up About a minute          client
ddca93d08801   lab-2-load_balancer "bash"                  About a minute ago Up About a minute          load_balancer
dbcacde7622e   lab-2-server2   "bash"                  9 days ago    Up About a minute          server2
25dc7e033916   lab-2-server1   "bash"                  9 days ago    Up About a minute          server1

[~/C0513/Lab-2]
kalindu ➤
```

## Task 1: Set Up Server Simulations

Logged into the containers for Server 1 and Server 2 and started basic Python HTTP servers.

```
[~/C0513/Lab-2]
kalindu ➤ sudo docker exec -it server1 bash
root@25dc7e033916:/# mkdir /web && cd /web
root@25dc7e033916:/web# echo "Hello from Servers!" > index.html
root@25dc7e033916:/web# python3 -m http.server 80
Serving HTTP on 0.0.0.0 port 80 (http://0.0.0.0:80/) ...
```

```
[~/C0513/Lab-2]
kalindu ➤ sudo docker exec -it server2 bash
root@dbcacde7622e:/# mkdir /web && cd /web
root@dbcacde7622e:/web# echo "Hello from Servers!" > index.html
root@dbcacde7622e:/web# python3 -m http.server 80
Serving HTTP on 0.0.0.0 port 80 (http://0.0.0.0:80/) ...
```

From the client container, both servers were accessed using curl for initial availability testing

```
[~/C0513/Lab-2]
kalindu ➤ sudo docker exec -it client bash
root@6c3da392631c:/# curl http://server1
Hello from Servers!
root@6c3da392631c:/# curl http://server2
Hello from Servers!
root@6c3da392631c:/#
```

Since an external client was used, a proxy was required to forward traffic from the client through the load balancer to the servers, thus creating the proxy.py script attached in [Appendix 3](#).

The following is a snapshot of logging into the load balancer Docker container and executing the proxy.py script.

```
[~/C0513/Lab-2]
kalindu ➤ sudo docker exec -it load_balancer bash
[sudo] password for kalindu:
root@ddca93d08801:/# sudo vim proxy.py
root@ddca93d08801:/# sudo python3 proxy.py
Proxy running on port 80...

=== Server Usage Stats ===
172.28.1.11: 0 requests
172.28.1.12: 0 requests
=====
```

The tool “wrk” was used for load testing by generating high traffic rates to simulate realistic traffic conditions. The following command was executed to send concurrent requests: “wrk -t4 -c50 -d30s <http://172.28.1.10>”, which ran a 30-second test with 4 threads and 50 connections against the load balancer at 172.28.1.10.

Initially, traffic was sent directly from the client without a load balancer configured, and the following outputs confirm that only a single server handled the requests.

Client-side traffic generation:

```
root@6c3da392631c:/# wrk -t4 -c50 -d30s http://172.28.1.10
Running 30s test @ http://172.28.1.10
 4 threads and 50 connections
  Thread Stats   Avg      Stdev     Max   +/-  Stdev
    Latency    23.54ms   111.58ms   1.67s   96.60%
    Req/Sec    316.68    192.53    1.03k   65.54%
 23783 requests in 30.06s, 6.31MB read
Socket errors: connect 0, read 0, write 0, timeout 6
Requests/sec:    791.26
Transfer/sec:    214.81KB
root@6c3da392631c:/#
```

Proxy-server logs:

```
172.28.1.20 - - [18/May/2025 13:54:13] "GET / HTTP/1.1" 200 -
172.28.1.20 - - [18/May/2025 13:54:13] "GET / HTTP/1.1" 200 -
172.28.1.20 - - [18/May/2025 13:54:13] "GET / HTTP/1.1" 200 -
172.28.1.20 - - [18/May/2025 13:54:13] "GET / HTTP/1.1" 200 -
172.28.1.20 - - [18/May/2025 13:54:13] "GET / HTTP/1.1" 200 -
172.28.1.20 - - [18/May/2025 13:54:13] "GET / HTTP/1.1" 200 -
172.28.1.20 - - [18/May/2025 13:54:13] "GET / HTTP/1.1" 200 -
172.28.1.20 - - [18/May/2025 13:54:13] "GET / HTTP/1.1" 200 -
172.28.1.20 - - [18/May/2025 13:54:13] "GET / HTTP/1.1" 200 -
172.28.1.20 - - [18/May/2025 13:54:13] "GET / HTTP/1.1" 200 -
172.28.1.20 - - [18/May/2025 13:54:13] "GET / HTTP/1.1" 200 -
172.28.1.20 - - [18/May/2025 13:54:13] "GET / HTTP/1.1" 200 -
172.28.1.20 - - [18/May/2025 13:54:13] "GET / HTTP/1.1" 200 -
172.28.1.20 - - [18/May/2025 13:54:13] "GET / HTTP/1.1" 200 -
172.28.1.20 - - [18/May/2025 13:54:13] "GET / HTTP/1.1" 200 -

=== Server Usage Stats ===
172.28.1.11: 0 requests
172.28.1.12: 23790 requests
=====
```

## Task 2: Create Load Balancer Script (Bash)

Then the `load_balancer.sh` script was created which is attached in [Appendix 4](#).

This script continuously simulates random load values for two servers and selects the server with the lower load by writing its IP to a file (`/tmp/selected_server`), updating the selection every 5 seconds; the selected server IP is then used by the proxy to forward traffic accordingly.

Load balancer logs: (Indicates simulated random loads and the selected server per each 5 seconds)

```
-----  
Simulated random server loads:  
Server 1 (172.28.1.11) load: 91  
Server 2 (172.28.1.12) load: 9  
Selected server: 172.28.1.12  
-----  
Simulated random server loads:  
Server 1 (172.28.1.11) load: 54  
Server 2 (172.28.1.12) load: 17  
Selected server: 172.28.1.12  
-----  
Simulated random server loads:  
Server 1 (172.28.1.11) load: 18  
Server 2 (172.28.1.12) load: 79  
Selected server: 172.28.1.11  
-----  
Simulated random server loads:  
Server 1 (172.28.1.11) load: 70  
Server 2 (172.28.1.12) load: 13  
Selected server: 172.28.1.12  
-----  
Simulated random server loads:  
Server 1 (172.28.1.11) load: 82  
Server 2 (172.28.1.12) load: 51  
Selected server: 172.28.1.12  
-----  
Simulated random server loads:  
Server 1 (172.28.1.11) load: 45  
Server 2 (172.28.1.12) load: 27  
Selected server: 172.28.1.12  
-----  
Simulated random server loads:  
Server 1 (172.28.1.11) load: 31  
Server 2 (172.28.1.12) load: 36  
Selected server: 172.28.1.11  
-----  
Simulated random server loads:  
Server 1 (172.28.1.11) load: 0  
Server 2 (172.28.1.12) load: 5  
Selected server: 172.28.1.11  
|
```

Client-side traffic generation:

```
root@6c3da392631c:/# wrk -t4 -c50 -d30s http://172.28.1.10
Running 30s test @ http://172.28.1.10
 4 threads and 50 connections
   Thread Stats   Avg      Stdev     Max   +/-  Stdev
   Latency    9.23ms   47.80ms   1.67s   99.11%
   Req/Sec   295.31   192.56   1.33k   68.91%
 11664 requests in 30.05s, 3.08MB read
Socket errors: connect 0, read 62, write 0, timeout 16
Non-2xx or 3xx responses: 62
Requests/sec:   388.12
Transfer/sec:   105.12KB
root@6c3da392631c:/#
```

Proxy-server logs:

```
172.28.1.20 - - [18/May/2025 13:55:41] "GET / HTTP/1.1" 200 -
172.28.1.20 - - [18/May/2025 13:55:41] "GET / HTTP/1.1" 200 -
172.28.1.20 - - [18/May/2025 13:55:41] "GET / HTTP/1.1" 200 -
172.28.1.20 - - [18/May/2025 13:55:41] "GET / HTTP/1.1" 200 -
172.28.1.20 - - [18/May/2025 13:55:41] "GET / HTTP/1.1" 200 -
172.28.1.20 - - [18/May/2025 13:55:41] "GET / HTTP/1.1" 200 -
172.28.1.20 - - [18/May/2025 13:55:41] "GET / HTTP/1.1" 200 -
172.28.1.20 - - [18/May/2025 13:55:41] "GET / HTTP/1.1" 200 -
172.28.1.20 - - [18/May/2025 13:55:41] "GET / HTTP/1.1" 200 -
172.28.1.20 - - [18/May/2025 13:55:41] "GET / HTTP/1.1" 200 -
172.28.1.20 - - [18/May/2025 13:55:41] "GET / HTTP/1.1" 200 -
172.28.1.20 - - [18/May/2025 13:55:41] "GET / HTTP/1.1" 200 -
172.28.1.20 - - [18/May/2025 13:55:41] "GET / HTTP/1.1" 200 -
172.28.1.20 - - [18/May/2025 13:55:41] "GET / HTTP/1.1" 200 -
172.28.1.20 - - [18/May/2025 13:55:41] "GET / HTTP/1.1" 200 -
172.28.1.20 - - [18/May/2025 13:55:41] "GET / HTTP/1.1" 200 -
172.28.1.20 - - [18/May/2025 13:55:41] "GET / HTTP/1.1" 200 -
172.28.1.20 - - [18/May/2025 13:55:41] "GET / HTTP/1.1" 200 -
172.28.1.20 - - [18/May/2025 13:55:41] "GET / HTTP/1.1" 200 -

=== Server Usage Stats ===
172.28.1.11: 6720 requests
172.28.1.12: 4888 requests
: 62 requests
=====
```

### Task 3: Enhance Logic Using Python

Finally, the `ai_loadbalancer.py` was created which is attached in [Appendix 5](#).

This Python script simulates a load balancer by selecting the server that has received fewer requests—tracked using counts derived from the proxy server logs—and writes the selected server’s IP to a file every second. This helps distribute traffic more evenly, and intelligently between the two servers over time.

Load balancer logs: (Indicates the selected server per each 1 second)

```
-----  
Selected server: 172.28.1.11  
-----  
-----  
Selected server: 172.28.1.12  
-----  
-----  
Selected server: 172.28.1.11  
-----  
-----  
Selected server: 172.28.1.12  
-----  
-----  
Selected server: 172.28.1.11  
-----  
-----  
Selected server: 172.28.1.12  
-----  
-----  
Selected server: 172.28.1.11  
-----  
-----  
Selected server: 172.28.1.12  
-----  
-----  
Selected server: 172.28.1.11  
-----  
-----  
Selected server: 172.28.1.12  
-----  
-----  
Selected server: 172.28.1.11  
-----  
-----  
Selected server: 172.28.1.12  
-----  
-----  
Selected server: 172.28.1.11  
-----
```

Client-side traffic generation:

```
root@6c3da392631c:/# wrk -t4 -c50 -d30s http://172.28.1.10
Running 30s test @ http://172.28.1.10
 4 threads and 50 connections
  Thread Stats   Avg      Stdev     Max   +/-  Stdev
    Latency    14.51ms   78.58ms   1.67s   98.01%
    Req/Sec    286.66    156.85   717.00    67.64%
 16912 requests in 30.04s, 4.48MB read
Socket errors: connect 0, read 9, write 0, timeout 14
Non-2xx or 3xx responses: 9
Requests/sec:    562.90
Transfer/sec:    152.78KB
root@6c3da392631c:/#
```

Proxy-server logs:

```
172.28.1.20 - - [18/May/2025 13:59:02] "GET / HTTP/1.1" 200 -
172.28.1.20 - - [18/May/2025 13:59:02] "GET / HTTP/1.1" 200 -
172.28.1.20 - - [18/May/2025 13:59:02] "GET / HTTP/1.1" 200 -
172.28.1.20 - - [18/May/2025 13:59:02] "GET / HTTP/1.1" 200 -
172.28.1.20 - - [18/May/2025 13:59:02] "GET / HTTP/1.1" 200 -
172.28.1.20 - - [18/May/2025 13:59:02] "GET / HTTP/1.1" 200 -
172.28.1.20 - - [18/May/2025 13:59:02] "GET / HTTP/1.1" 200 -
172.28.1.20 - - [18/May/2025 13:59:02] "GET / HTTP/1.1" 200 -
172.28.1.20 - - [18/May/2025 13:59:02] "GET / HTTP/1.1" 200 -
172.28.1.20 - - [18/May/2025 13:59:02] "GET / HTTP/1.1" 200 -
172.28.1.20 - - [18/May/2025 13:59:04] "GET / HTTP/1.1" 200 -
172.28.1.20 - - [18/May/2025 13:59:04] "GET / HTTP/1.1" 200 -
172.28.1.20 - - [18/May/2025 13:59:04] "GET / HTTP/1.1" 200 -
172.28.1.20 - - [18/May/2025 13:59:04] "GET / HTTP/1.1" 200 -
172.28.1.20 - - [18/May/2025 13:59:04] "GET / HTTP/1.1" 200 -
172.28.1.20 - - [18/May/2025 13:59:04] "GET / HTTP/1.1" 200 -

=== Server Usage Stats ===
172.28.1.11: 8262 requests
172.28.1.12: 8651 requests
: 9 requests
=====
```

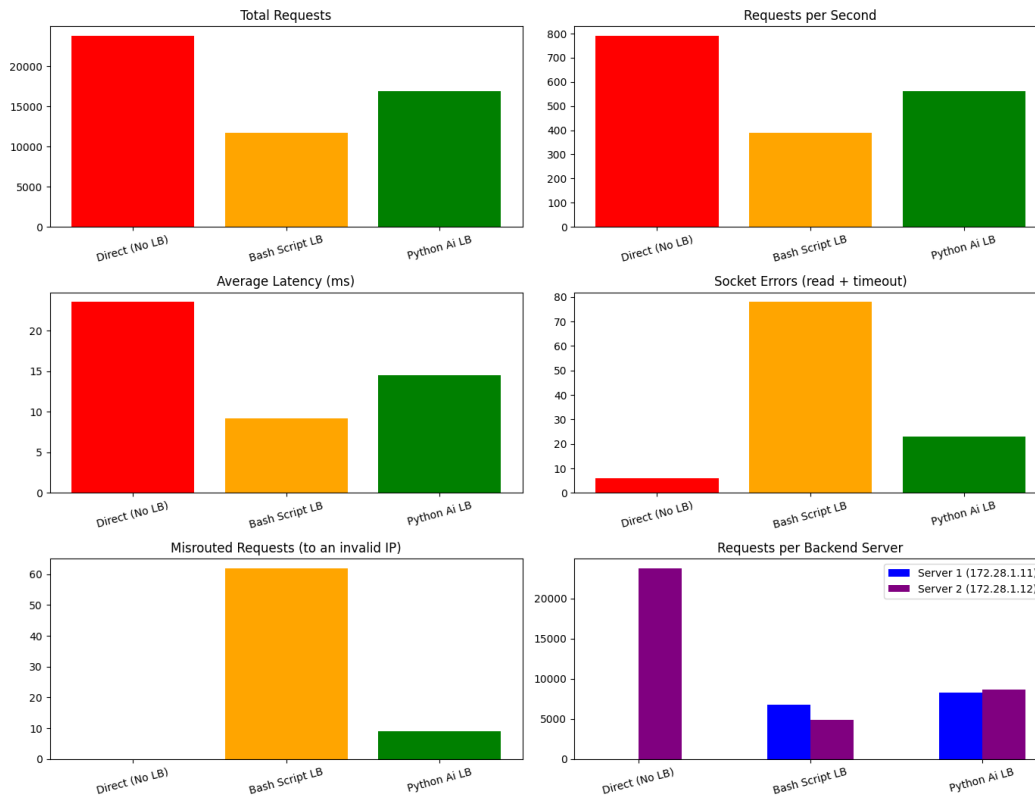


## Task 4: Analyze Performance and Decisions

The following is the comparison between

- No Load Balancer
- Bash script Load Balancer (Based on random loads)
- Python scrip Load Balancer (Based on history logs and intelligent selection)

Load Balancer Performance Comparison



Metric	Best Performer	Comments
Throughput	No LB > AI LB > Bash LB	Requests were consistently routed to a single server, avoiding selection overhead in No LB.
Latency	Bash LB > AI LB > No LB	Random selection led to more distributed load early, reducing queuing and keeping latency low in the Bash LB.
Reliability	No LB > AI LB > Bash LB	Fewer rerouting decisions and more consistent routing reduced response failures in No LB.
Distribution	AI LB > Bash LB >>> No LB	AI LB tracked usage and adjusted accordingly, ensuring nearly even server utilization.



Comparing the two Load Balancer implementations;

- AI LB handled load smarter, leading to better throughput and more even server distribution.
- Bash LB had slightly lower latency, likely due to early random load spreading.
- AI LB was more reliable, with fewer errors and failed responses.
- Overall, AI LB scaled better and is more adaptable for real-world use.

## Appendix-1: docker-compose.yaml

```
version: '3.8'

services:
  load_balancer:
    container_name: load_balancer
    networks:
      mynetwork:
        ipv4_address: 172.28.1.10
    tty: true
    stdin_open: true
    command: bash
    privileged: true
    build:
      context: .
      dockerfile: Dockerfile

  server1:
    container_name: server1
    networks:
      mynetwork:
        ipv4_address: 172.28.1.11
    tty: true
    stdin_open: true
    command: bash
    privileged: true
    build:
      context: .
      dockerfile: Dockerfile

  server2:
    container_name: server2
    networks:
      mynetwork:
        ipv4_address: 172.28.1.12
    tty: true
    stdin_open: true
    command: bash
    privileged: true
    build:
      context: .
      dockerfile: Dockerfile

  client:
```

```
container_name: client
networks:
  mynetwork:
    ipv4_address: 172.28.1.20
tty: true
stdin_open: true
command: bash
privileged: true
build:
  context: .
  dockerfile: Dockerfile

networks:
  mynetwork:
    driver: bridge
    ipam:
      config:
        - subnet: 172.28.1.0/24
```

## Appendix-2: Dockerfile used for all the nodes (servers, clients, load balancer)

```
FROM ubuntu:20.04

ENV DEBIAN_FRONTEND=noninteractive

RUN apt-get update && apt-get install -y \
    python3 \
    iperf3 \
    iproute2 \
    net-tools \
    iputils-ping \
    bash \
    sudo \
    build-essential \
    libssl-dev \
    git \
    unzip \
    && apt-get clean

# Install wrk from source
RUN git clone https://github.com/wg/wrk.git /opt/wrk && \
    make -C /opt/wrk && \
    ln -s /opt/wrk/wrk /usr/local/bin/wrk
```

### Appendix-3: proxy.py

```
from http.server import BaseHTTPRequestHandler, HTTPServer
import requests
import time
import threading

SELECTED_FILE = "/tmp/selected_server"
PORT = 80

# In-memory usage counter
request_count = {
    "172.28.1.11": 0,
    "172.28.1.12": 0
}

class ProxyHandler(BaseHTTPRequestHandler):
    def do_GET(self):
        try:
            # Read selected server IP
            with open(SELECTED_FILE, 'r') as f:
                backend = f.read().strip()

            # Track request count
            if backend in request_count:
                request_count[backend] += 1
            else:
                request_count[backend] = 1 # in case new IP appears

            url = f"http://{backend}{self.path}"

            # Forward request to server
            resp = requests.get(url, timeout=2)

            # Send response headers
            self.send_response(resp.status_code)
            for k, v in resp.headers.items():
                if k.lower() != 'transfer-encoding':
                    self.send_header(k, v)
            self.end_headers()

            # Send response content
            try:
                self.wfile.write(resp.content)
            except BrokenPipeError:
```

```

        pass

    except Exception as e:
        self.send_response(502)
        self.end_headers()
        try:
            self.wfile.write(f"Error: {e}".encode())
        except BrokenPipeError:
            pass

def log_stats():
    while True:
        time.sleep(5)
        print("\n=== Server Usage Stats ===")
        for server, count in request_count.items():
            print(f"{server}: {count} requests")
        print("=====\n")

if __name__ == "__main__":
    print(f"Proxy running on port {PORT}...")

    # Start background logging thread
    stats_thread = threading.Thread(target=log_stats, daemon=True)
    stats_thread.start()

    # Start HTTP proxy server
    server = HTTPServer(('', PORT), ProxyHandler)
    server.serve_forever()

```

#### Appendix-4: load\_balancer.sh

```
#!/bin/bash

SERVER1="172.28.1.11"
SERVER2="172.28.1.12"
PORT="80"
SELECTED_FILE="/tmp/selected_server"

# Initial target
echo "$SERVER1" > "$SELECTED_FILE"

while true; do
    # Simulate random load values between 0 and 100
    LOAD1=$(( RANDOM % 101 ))
    LOAD2=$(( RANDOM % 101 ))

    echo "-----"
    echo "Simulated random server loads:"
    echo "Server 1 ($SERVER1) load: $LOAD1"
    echo "Server 2 ($SERVER2) load: $LOAD2"

    if [ "$LOAD1" -le "$LOAD2" ]; then
        TARGET="$SERVER1"
    else
        TARGET="$SERVER2"
    fi

    echo "$TARGET" > "$SELECTED_FILE"
    echo "Selected server: $TARGET"

    sleep 5
done
```



## Appendix-5: ai\_load\_balancer.py

```
import time

SERVER1 = "172.28.1.11"
SERVER2 = "172.28.1.12"
SELECTED_FILE = "/tmp/selected_server"

# Track how many times each server was selected
request_history = {
    SERVER1: 0,
    SERVER2: 0
}

def select_server():
    # Pick the server with fewer requests
    if request_history[SERVER1] <= request_history[SERVER2]:
        return SERVER1
    else:
        return SERVER2

def write_selected_server(server):
    with open(SELECTED_FILE, "w") as f:
        f.write(server)

def main():
    print("AI Load Balancer started using request history tracking...\n")

    while True:
        selected = select_server()
        request_history[selected] += 1
        write_selected_server(selected)

        print("-----")
        print(f"Selected server: {selected}")
        print("-----")

        time.sleep(1)

if __name__ == "__main__":
    main()
```